

# ML project review checklist

MUST-ASK QUESTIONS IN BOLD

## High-level questions about the project

- **What question were you trying to answer? How did you frame it as an ML task?**
- What is human-level performance on that task? What level of performance is needed?
- Is it possible to approach this problem without machine learning?
- If the analysis focused on deep learning methods, did you try shallow learning methods?
- What are the ethical and legal aspects of this project?
- **Which domain experts were involved in this analysis?**
- Which data scientists were involved in this analysis?
- Which tools or framework did you use? (How much of a known quantity is it?)
- Where is the pipeline published? (E.g. public or internal git repositories.)
- How thorough is the documentation?

## Questions about the data preparation

- Where did the feature data come from?
- Where did the labels come from?
- What kind of data exploration did you do?
- How did you clean the data? How long did this take?
- **Are the classes balanced? How did the distribution change your workflow?**
- What kind of normalization did you do?
- What did you do about missing data? E.g. what kind of imputation did you do?
- What kind of feature engineering did you do?
- How did you split the data into train, validate and test?

## Questions about training and evaluation

- Which models did you explore and why? Did you also try the simplest models that fit the problem?
- How did you tune the hyperparameters of the model? Did you try grid search or other methods?
- **What kind of validation did you do? Did you use cross-validation? How did you choose the folds?**
- What evaluation metric are you using? Why is it the most appropriate one?
- **How do training, validation, and test metrics compare?**
- If this was a classification task, how does a dummy classifier score?
- How are errors/residuals distributed? (Ideally normally distributed and homoscedastic.)
- How interpretable is your model? That is, do the learned parameters mean anything, and can we learn from them? E.g. what is the feature importance?
- If this was a classification task, are probabilities available in your model and did you use them?
- If this was a regression task, have you checked the residuals for normality and homoscedasticity?
- Are there benchmarks for this task, and how well does your model do on them?

## Next steps for the project

- How will you improve the model?
- Would collecting more data help? Can we address the imbalance with more data?
- Are there human or computing resources you need access to?
- How will you deploy the model?